Lesson 03 Demo 06

Autoscaling and Load Balancing with EKS

Objective: To implement containerization of applications on an AWS EKS cluster, including autoscaling and load balancing for improved scalability and availability

Tools required: AWS Management Console

Prerequisites: A ready EKS cluster with a worker node group

Steps to be followed:

1. Create AWS EKS autoscaling

Step 1: Create AWS EKS autoscaling

1.1 Sign in to the AWS Management Console

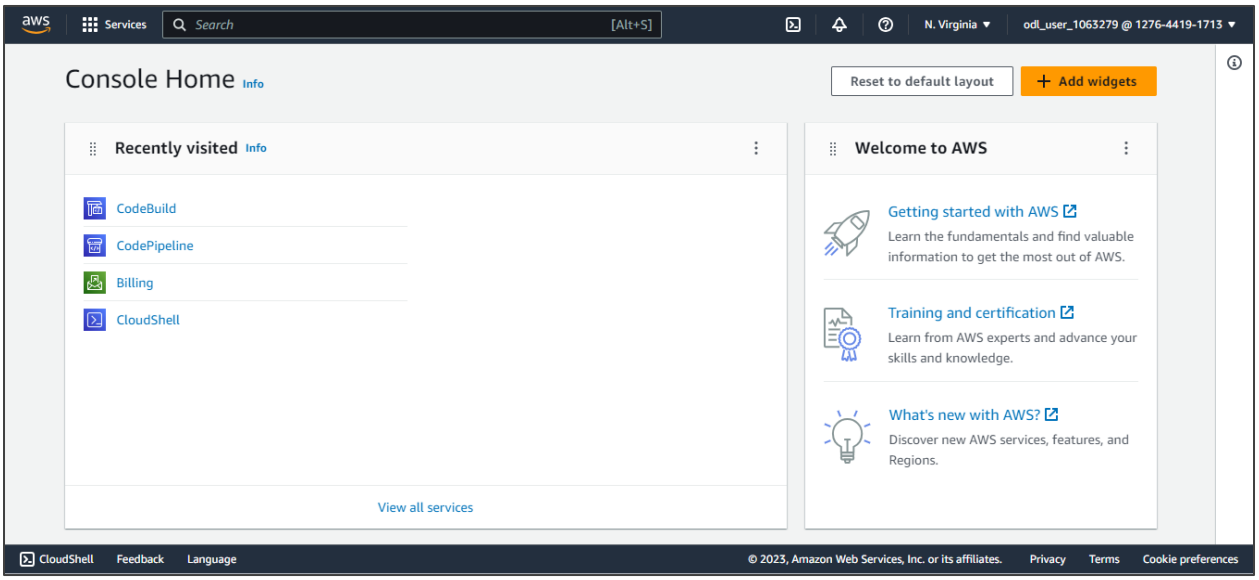

1.2 Search for and click on **Elastic Kubernetes Service**

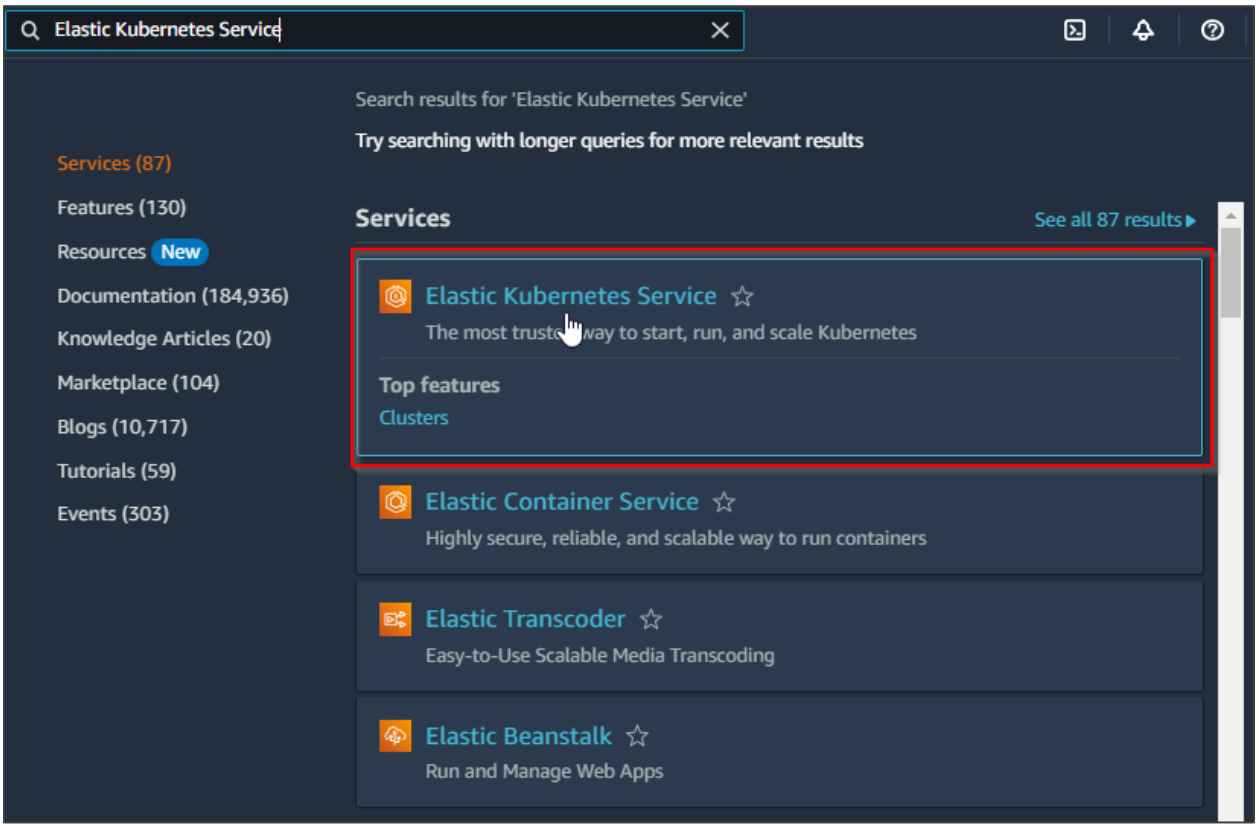

1.3 Select the desired EKS cluster

Note: Refer to Demo 04 of Lesson 06 to create the EKS cluster and worker node group

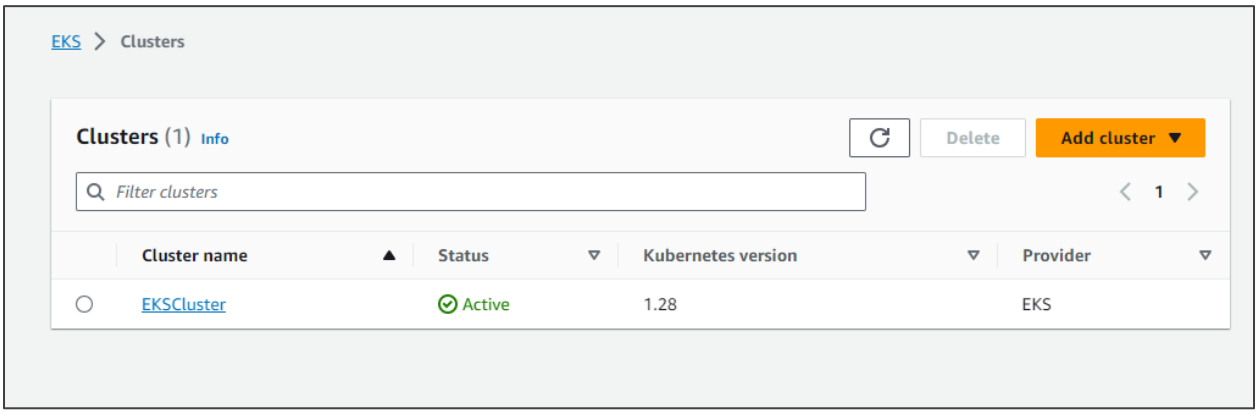

1.4 Click on **Compute**

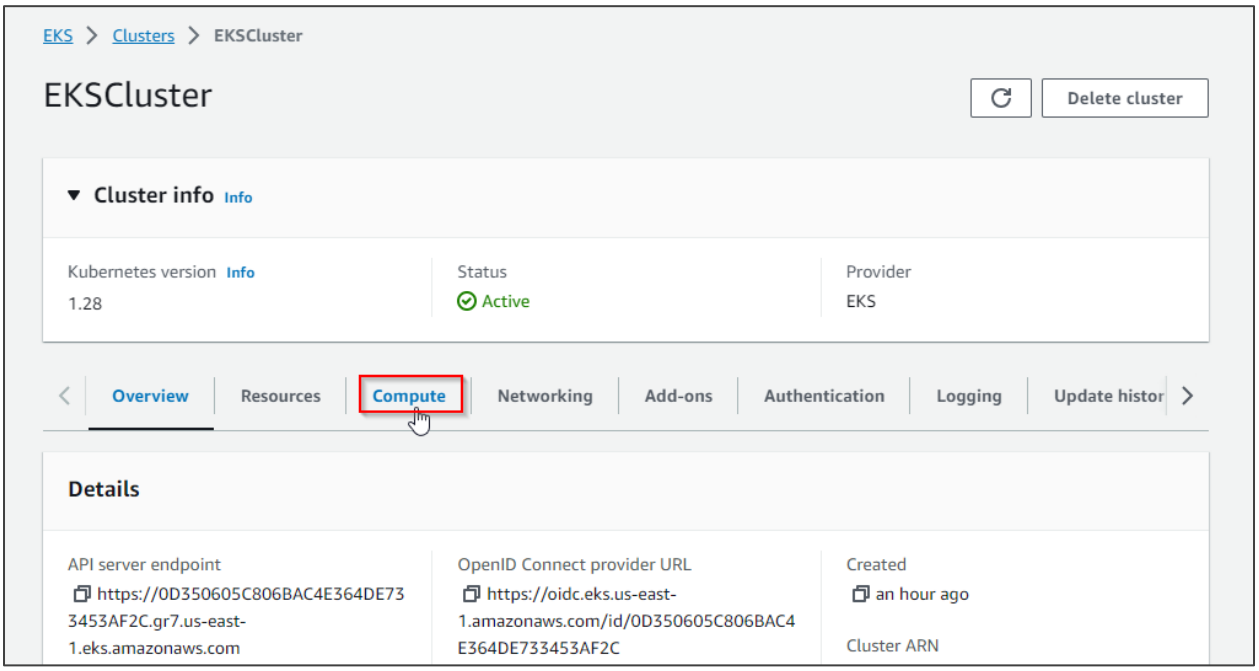

1.5 Click on the desired node group

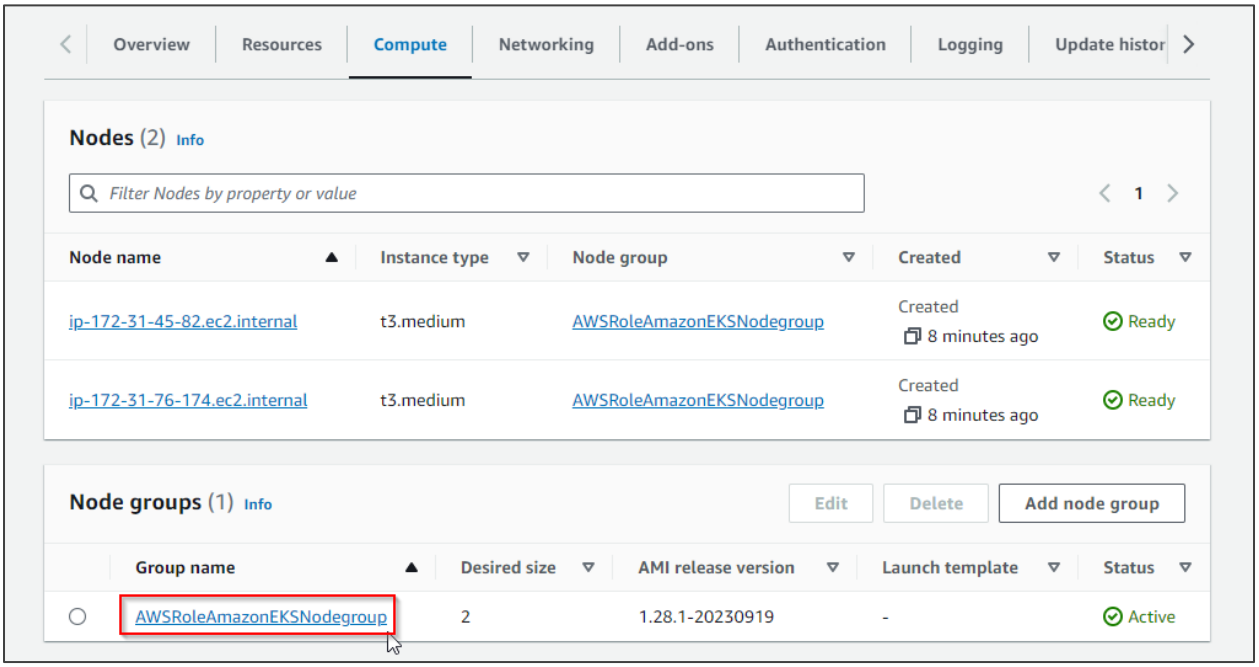

1.6 In the **Details** section, click on the **Autoscaling group name** link; a new window will appear.

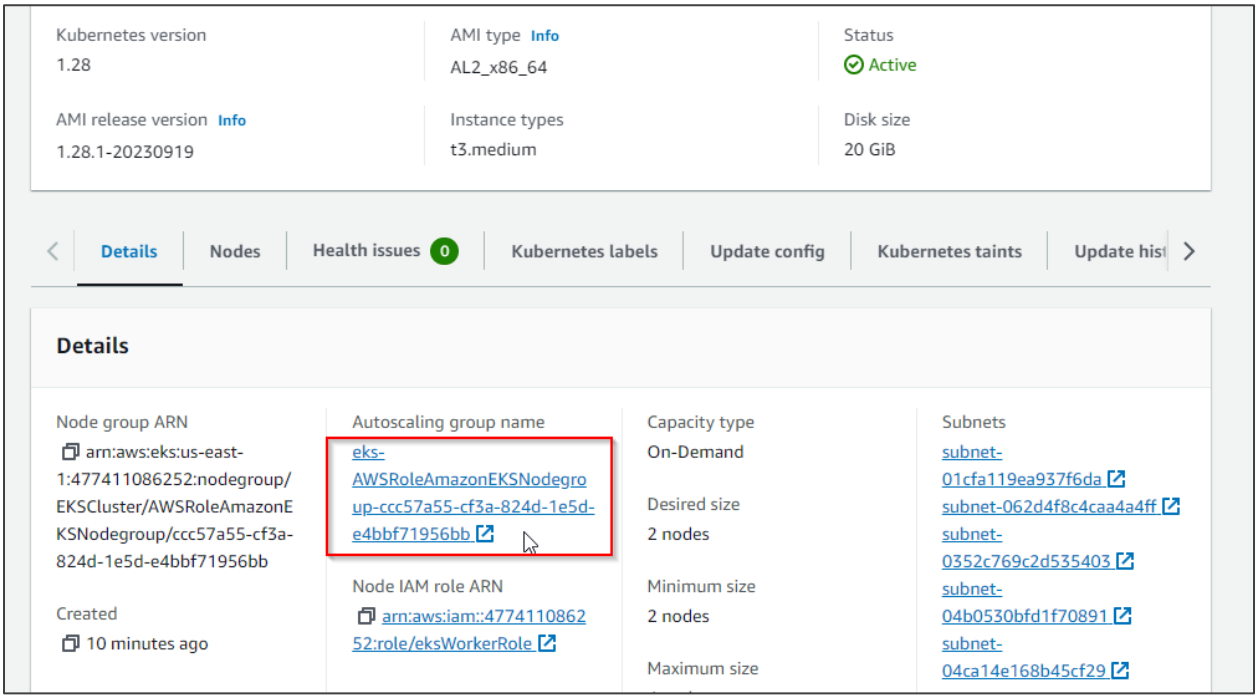

1.7 In the new window, click the **Edit** button

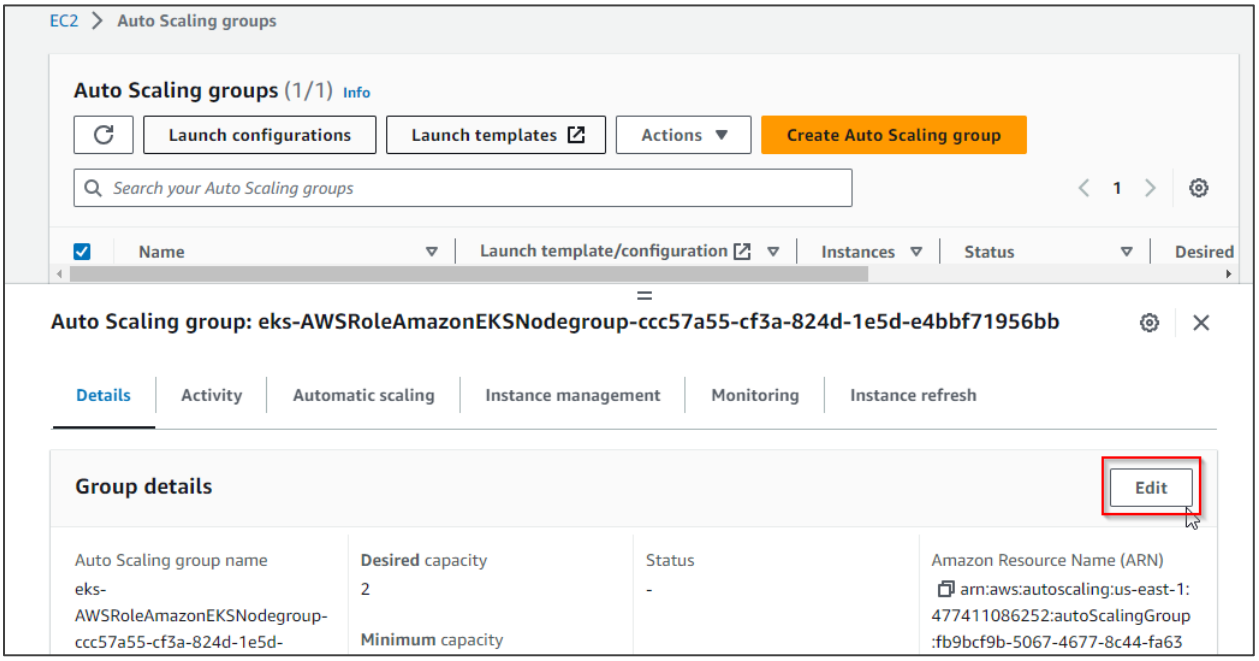

1.8 Add the desired capacity, minimum capacity, and maximum capacity, and then click the **Update** button

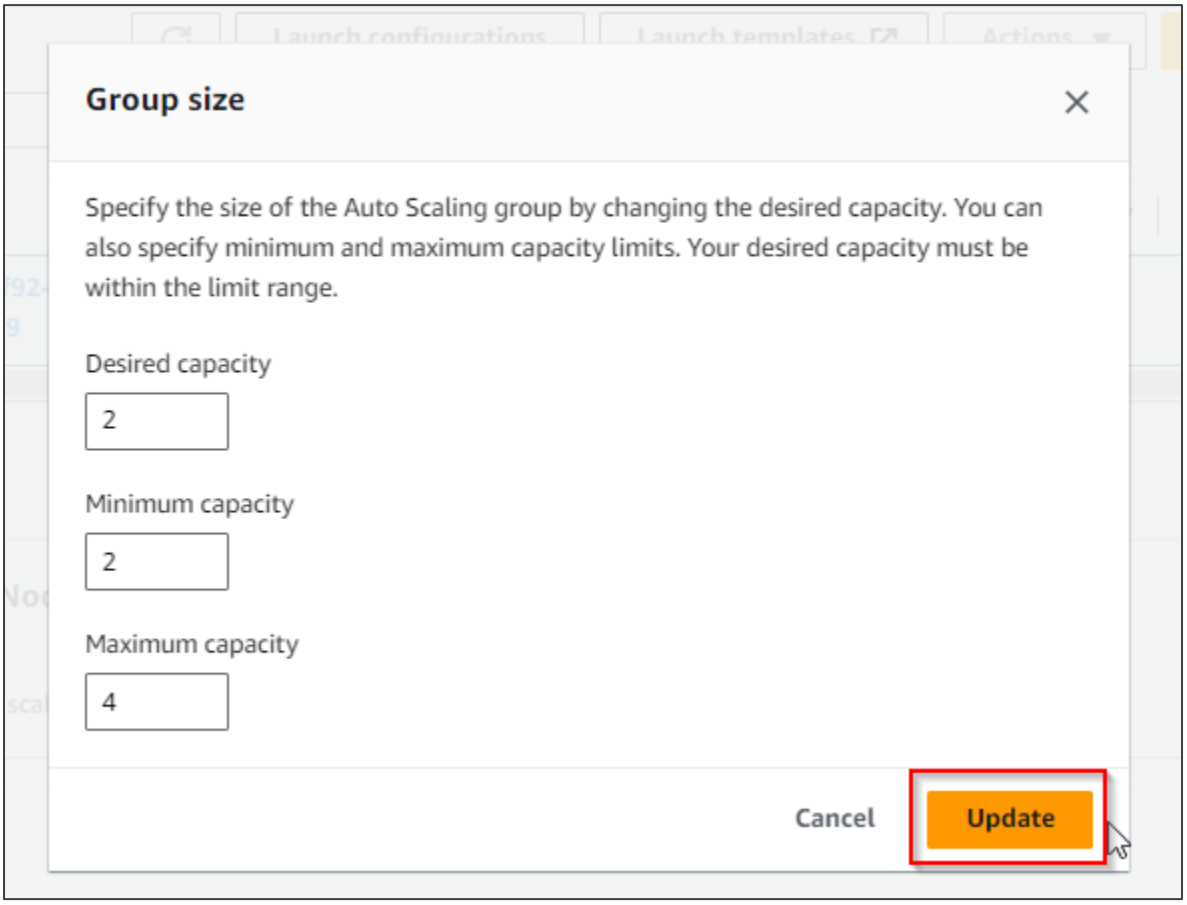

1.9 Click on **Automatic scaling**

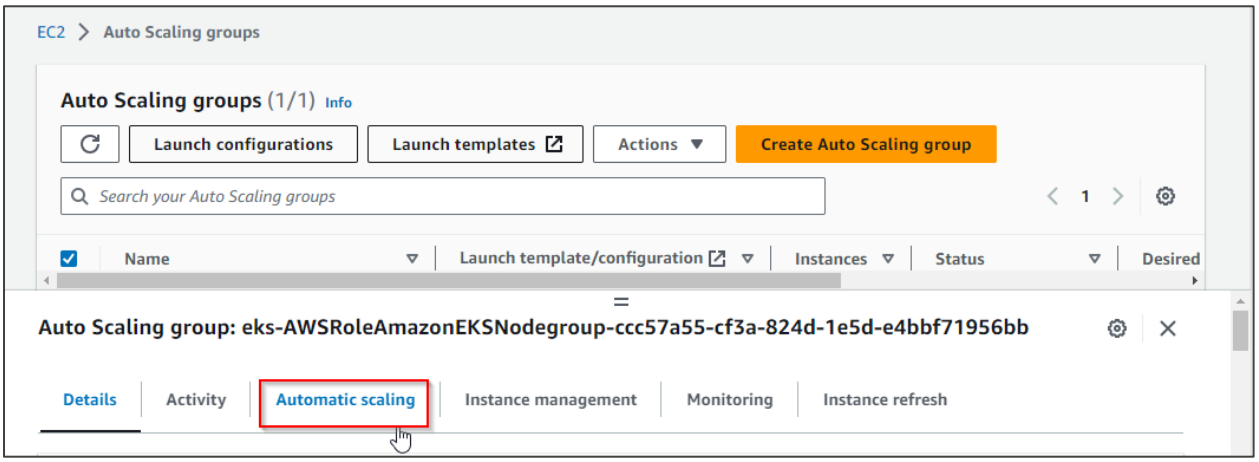

1.10 Click on the **Create dynamic scaling policy** button

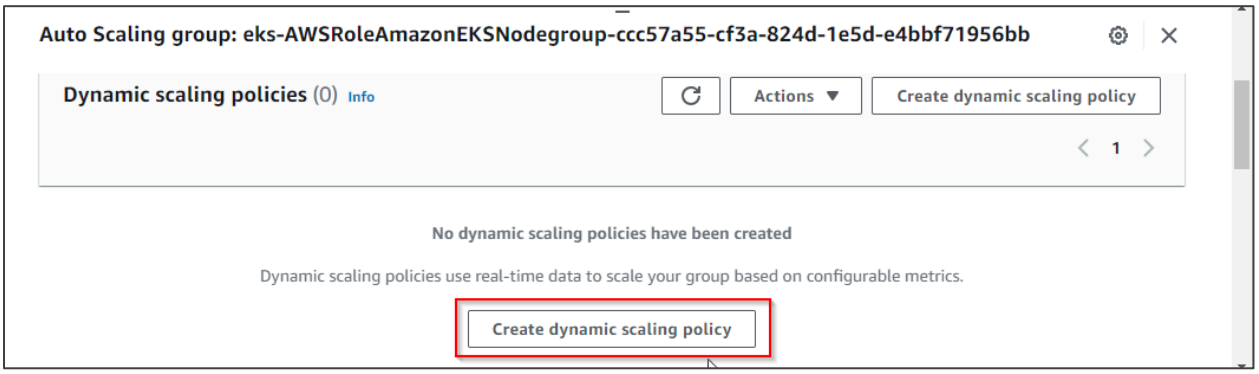

1.11 Enter the details to dynamically allocate resources to adjust for increases or decreases in traffic, and then click the **Create** button as shown in the screenshot below:

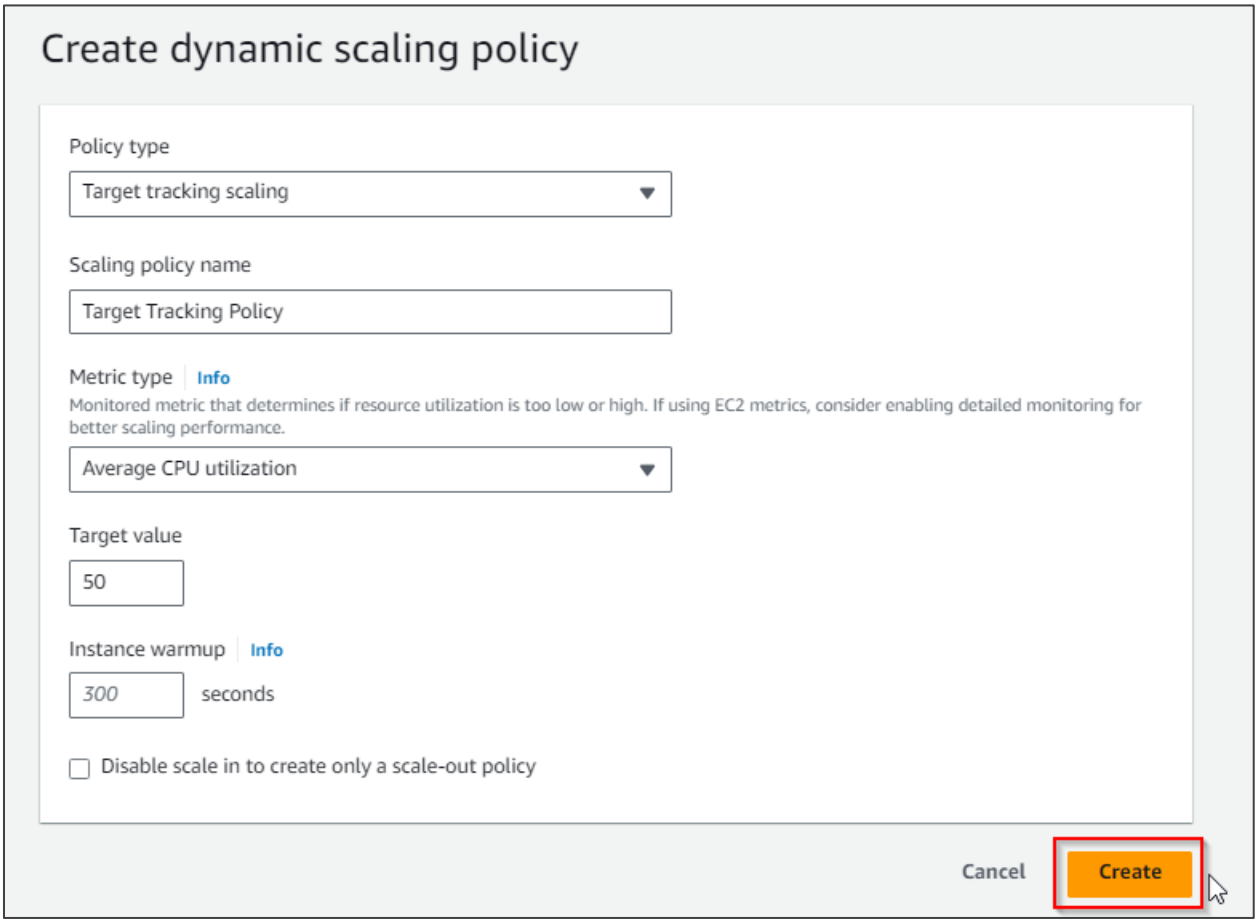

You will see the following interface:

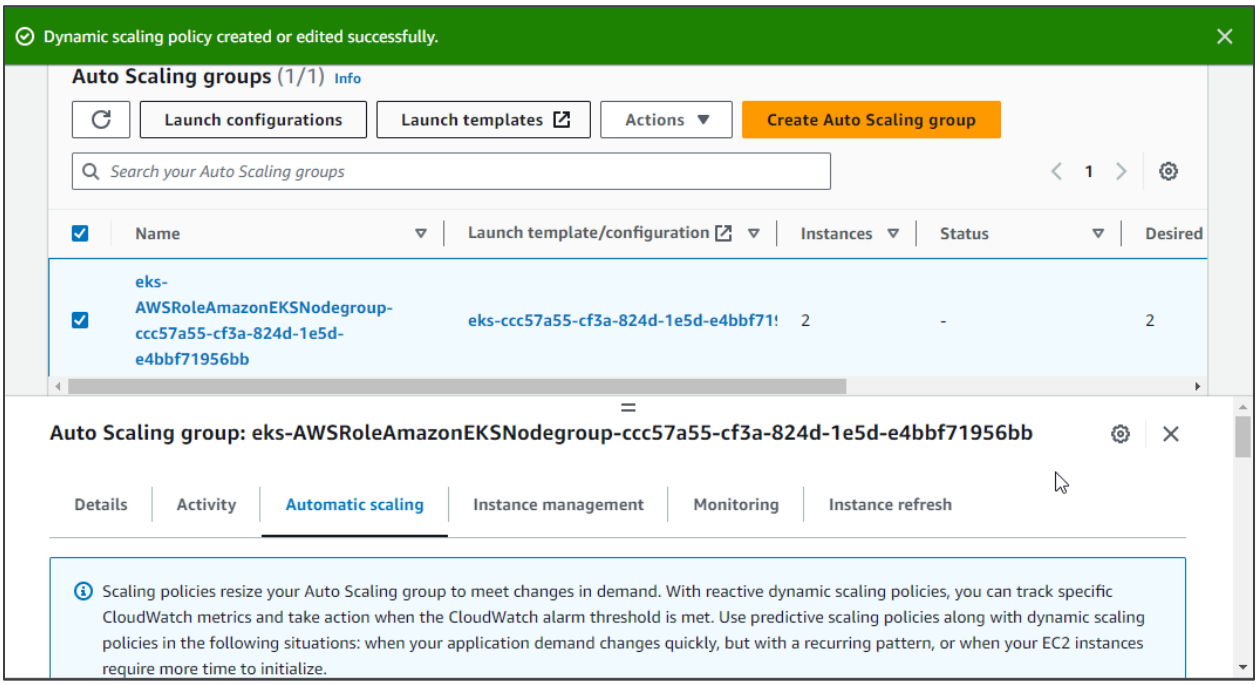

1.12 Navigate back to the previous window and click on **EKSCluster**

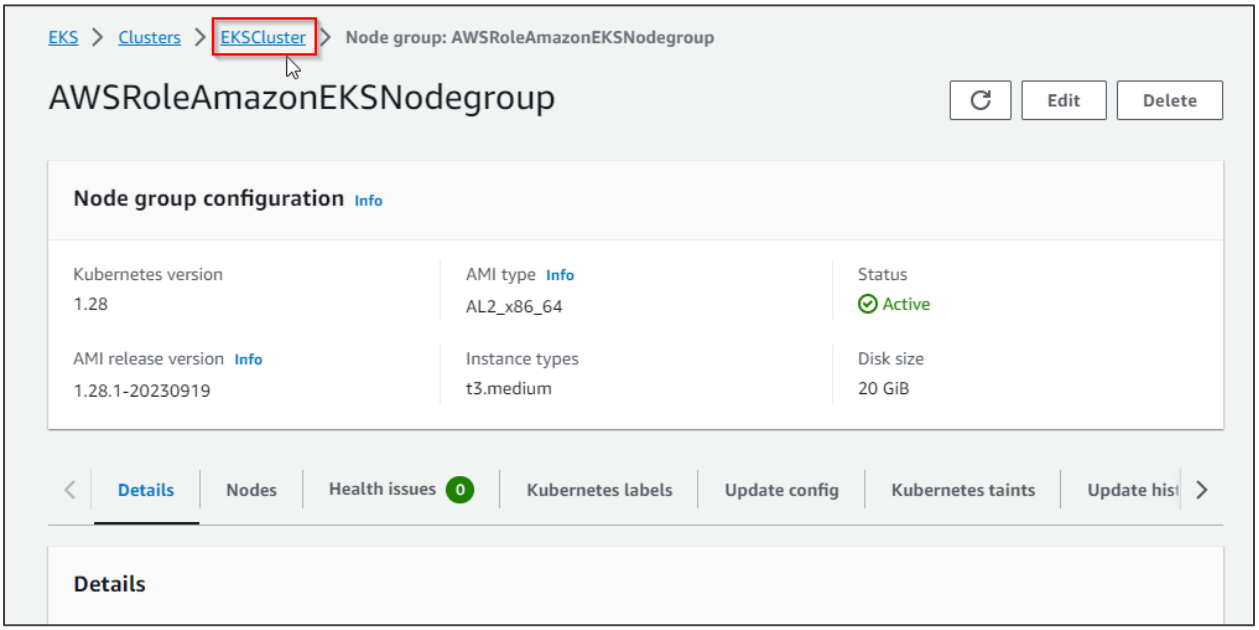

1.13 Click on **Resources**

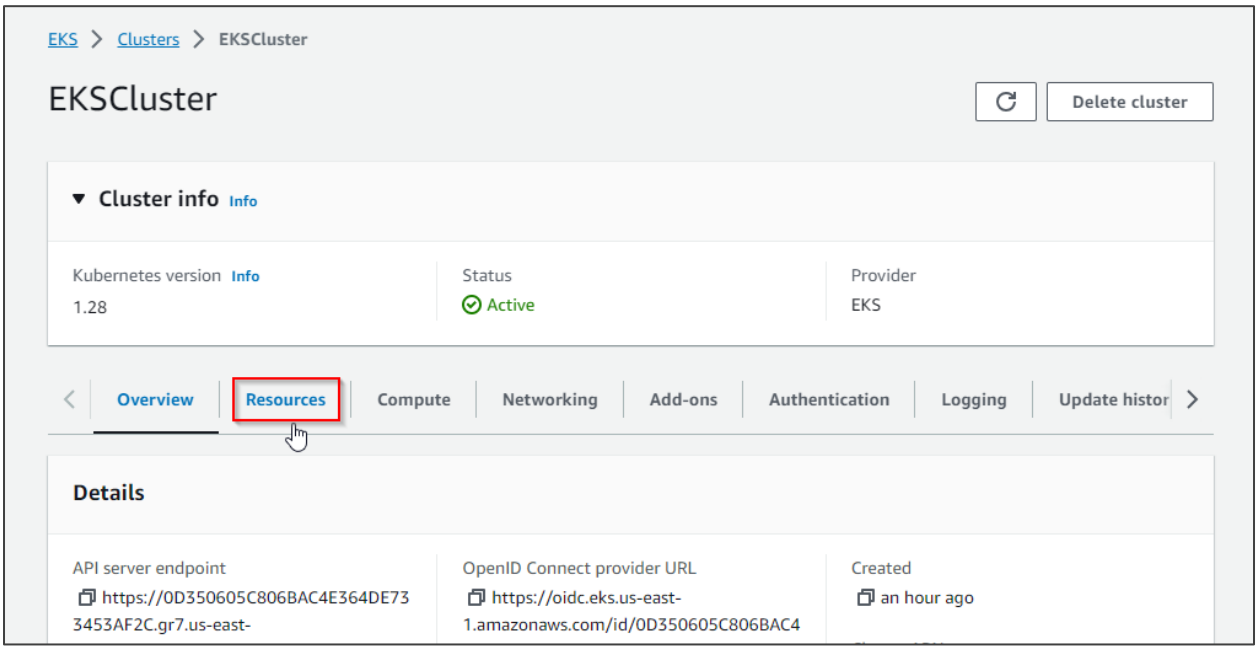

1.14 Click on **Pods** and select **kube-system** from the **All Namespaces** list

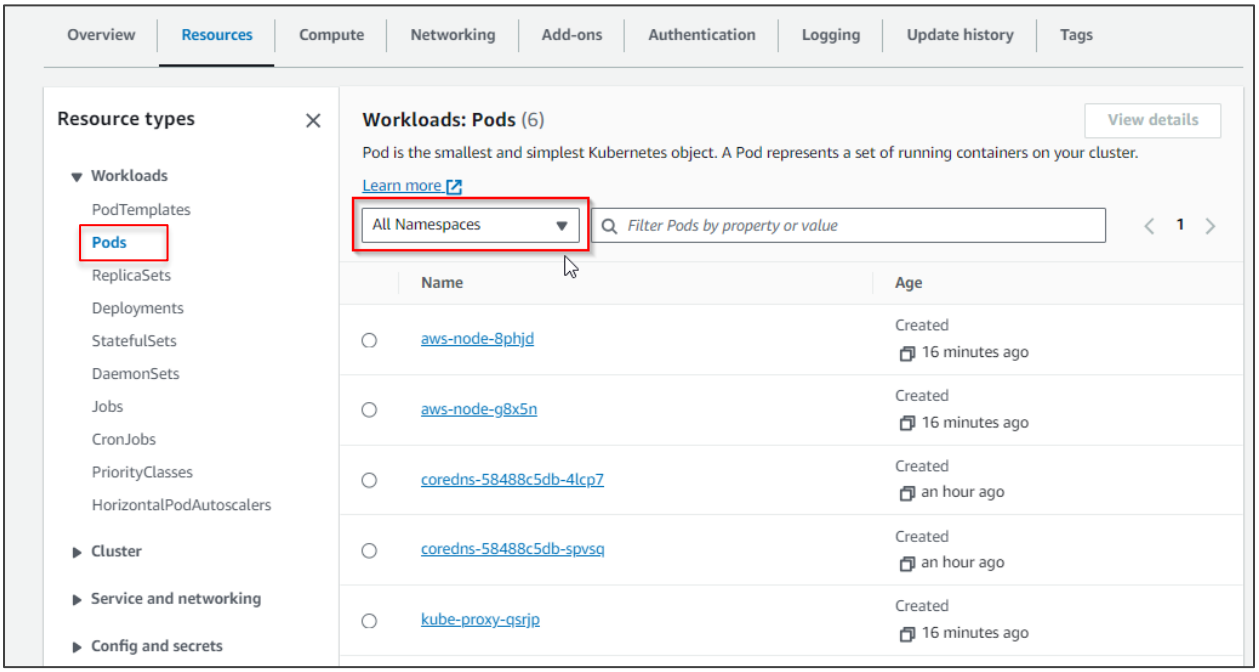

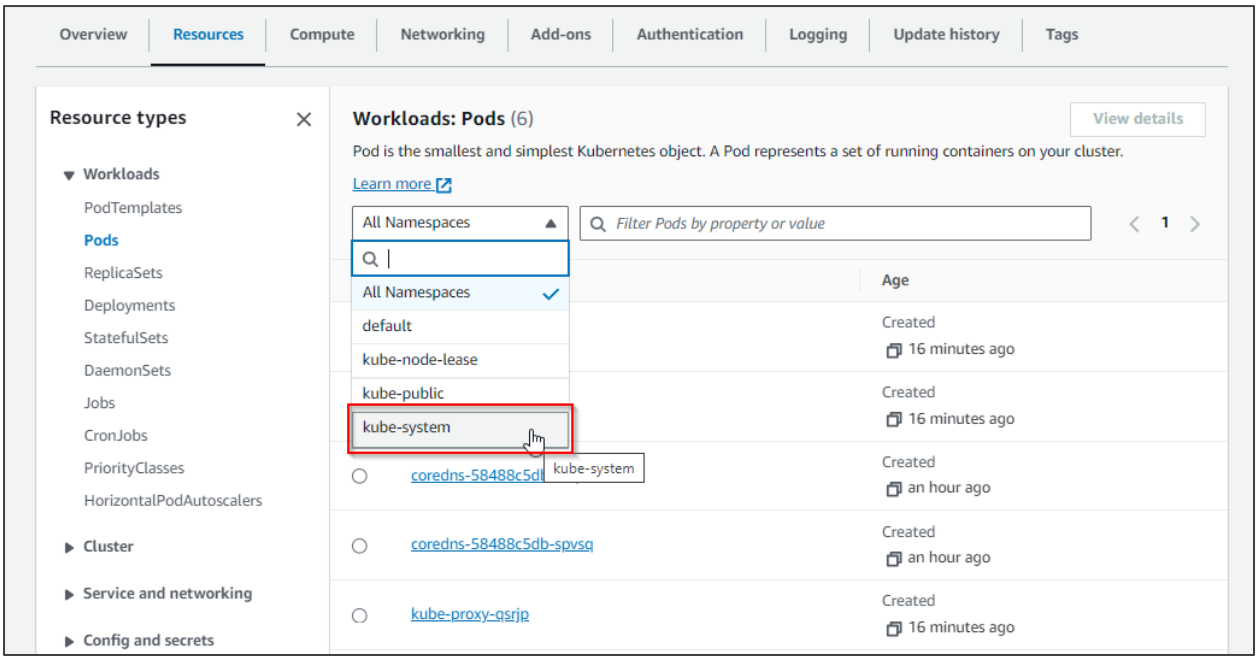

1.15 Select the desired workload name to validate the changes

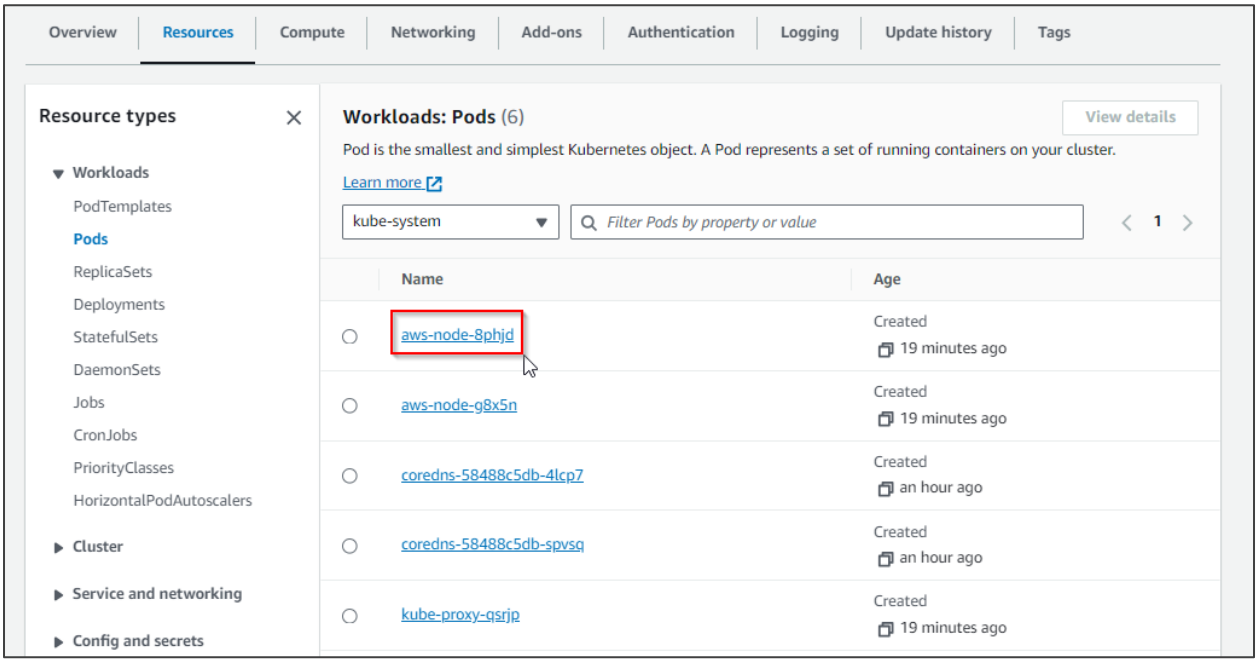

1.16 Click on the **aws-eks-nodeagent** container

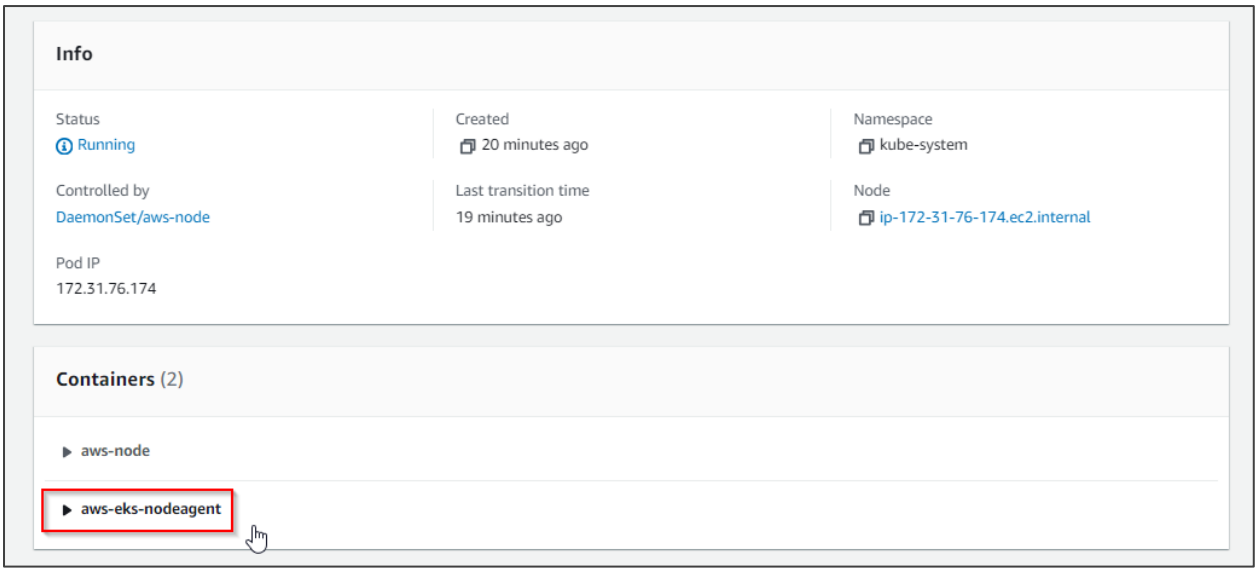

You will see the interface shown below:

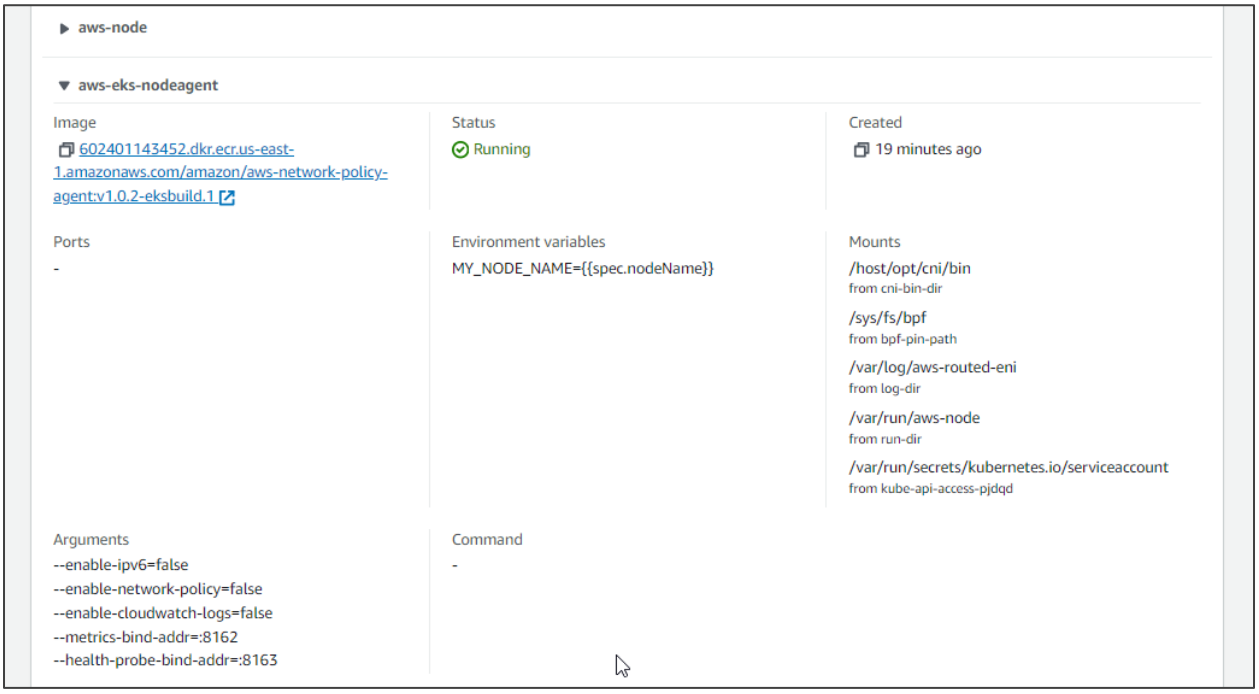

By following these steps, you have successfully implemented the containerization of applications on an AWS EKS cluster, including autoscaling and load balancing for improved scalability and availability.